

# CSC 555 Mining Big Data

## Assignment 3

**Due Tuesday, February 11<sup>th</sup>**

1) MapReduce:

a) Describe how to implement the following query in MapReduce

```
SELECT SUM(lo_extendedprice)
FROM lineorder, dwdate
WHERE lo_orderdate = d_datekey
      AND d_yearmonth = 'Jan1995'
      AND lo_discount BETWEEN 4 AND 7;
```

**Map lineorder : { lo\_orderdate: lo\_discount, lo\_extendedprice }**

**Map dwdate : { d\_datekey : d\_yearmonth }**

**WHERE lo\_orderdate = d\_datekey:::Join::: Reducer{ lineorder:  
d\_yearmonth, lo\_discount, lo\_extendedprice }**

**Map 3: { d\_yearmonth: lo\_discount, lo\_extendedprice }**

**Reducer d\_yearmonth = 'Jan1995' ..... Reducer{ d\_yearmonth = 'Jan1995' :  
lo\_discount, lo\_extendedprice }**

**Map 4: { d\_yearmonth , lo\_discount : lo\_extendedprice }**

**Reducer d\_yearmonth = 'Jan1995' ..... Reducer{ d\_yearmonth = 'Jan1995' ,  
lo\_discount BETWEEN 4 AND 7 : SUM(lo\_extendedprice)}**

b) **SELECT d\_month, COUNT(d\_sellingseason)**

```
FROM dwdate
GROUP BY d_month
ORDER BY COUNT(d_sellingseason)
```

**Mapper 1 dwdate: { d\_month : d\_sellingseason }**

**GROUP BY d\_month::: Mapper 2: { d\_month: d\_sellingseason }**

**Reduce: (key: d\_month, value: COUNT(d\_sellingseason))**

- 2) Consider a Hadoop job that processes an input data file of size equal to 88 disk blocks (88 different blocks, you can assume that HDFS replication factor is set to 1). The mapper in this job requires 2 minutes to read and fully process a single block of data. Reducer requires 1 second (**not** minute) to produce an answer for one key worth of values and there are a total of 6000 **distinct** keys (mappers generate a lot of key-value pairs, but keys only occur in the 1-6000 range for a total of 6000 unique entries). Assume that each node has a reducer and that the keys are distributed evenly.

- a) How long will it take to complete the job if you only had one Hadoop worker node? For the sake of simplicity, assume that that only one mapper and only one reducer are created on every node.

**The only mapper has to go to every block to do mapping so its  $88 \times 2 = 176$  mins, then the only one reducer so it's 6000 sec cost equal to 100 mins total is 276 mins.**

- b) 30 Hadoop worker nodes?

**30 nodes workers, 88 blocks divide 30 nodes is 2.93. In no any failure situation, workers have to do three rounds so is  $2 \times 3 = 6$  mins. For 30 reducers, which is  $6000/30 \text{ sec} = 200 \text{ sec}$  is equal to 3 mins and 20 sec. To sum up total time spend will be 9 min 20 sec.**

- c) 50 Hadoop worker nodes?

**50 nodes workers, 88 blocks divide 50 nodes is 1.76. In no any failure situation, workers have to do twos rounds so is  $2 \times 2 = 4$  mins. For 50 reducers, which is  $6000/50 \text{ sec} = 120 \text{ sec}$  is equal to 2 mins. To sum up total time spend will be 6 min.**

- d) 100 Hadoop worker nodes?

**100 nodes workers, 88 blocks divide 100 nodes is 0.88. In no any failure situation, workers have to do one rounds so is  $2 \times 1 = 2$  mins. For 100 reducers, which is  $6000/100 \text{ sec} = 60 \text{ sec}$  is equal to 1 min. To sum up total time spend will be 3 min.**

- e) Would changing the replication factor have any affect your answers for a-d?

**No, because the replication factor just makes sure if the node failure still keep the data and can sign worker into other nodes.**

You can ignore the network transfer costs as well as the possibility of node failure.

3)

- a) Suppose you have a 7-node cluster with replication factor of 3. Describe what MapReduce has to do after it determines that a node has crashed while a job is being processed. For simplicity, assume that the failed node is not replaced and your cluster is reduced to 6 nodes. Specifically:

- i) What does HDFS (the storage layer) have to do in response to node failure in this case?

**If name node doesn't receive any response from data node, then the name node, which is master node will replicate again into 6 nodes.**

- ii) What does MapReduce engine have to do to respond to the node failure? Assume that there was a job in progress because otherwise MapReduce does not need to do anything.

**Any Mapper or the reducer are completed their task, they will reset back to idle and eligible for schedule to other node that has the block.**

- b) Where does the Mapper store output key-value pairs before they are sent to Reducers?

**Mapper stored the data on local disks, or we called Intermediate files.**

- c) Can Reducers begin processing before Mapper phase is complete? **Why or why not?**

**The reducer have 3 phase, shuffle, sort, reduce so for shuffle, which is collecting data from mapper so this part might start when other node is working, but for sort, and reduce, reducer will only start when all the mapper done.**

- 4) Using the SSBM schema  
([http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM\\_schema\\_hive.sql](http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM_schema_hive.sql)) load the Part table into Hive (data available at <http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/part.tbl>)

**NOTE:** The schema above is made for Hive, but by default Hive assumes '\t' separated content. You will need to modify your CREATE TABLE statement to account for '|' delimiter in the data.

```
stop-dfs.sh: stop-yarn.sh: mr-jobhistory-daemon.sh stop historyserver  
rm -rf /tmp/hadoop-ec2-user/dfs/
```

```
hdfs namenode -format
```

```
start-dfs.sh: start-yarn.sh: mr-jobhistory-daemon.sh start historyserver
```

```
${HADOOP_HOME}/bin/hadoop fs -chmod g+w /tmp
```

```
${HADOOP_HOME}/bin/hadoop fs -chmod g+w /user/hive/warehouse
```

```
cd $HIVE_HOME
```

```
bin/hive
```

```
create table part (  
  p_partkey int,  
  p_name varchar(22),  
  p_mfgr varchar(6),  
  p_category varchar(7),  
  p_brand1 varchar(9),  
  p_color varchar(11),  
  p_type varchar(25),  
  p_size int,  
  p_container varchar(10)  
)
```

**ROW FORMAT DELIMITED FIELDS**

**TERMINATED BY '\t' STORED AS TEXTFILE:**

**LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl':**

**OVERWRITE INTO TABLE part:**

**select \* from part:**

**LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl'**

**OVERWRITE INTO TABLE part:**

**select \* from part:**

```
hive> select * from part;  
OK  
1      lace spring      MFGR#1 MFGR#11 MFGR#1121      goldenrod      PROMO BURNISHED COPPER 7      JUMBO PKG  
2      rosy metallic    MFGR#4 MFGR#43 MFGR#4318      blush      LARGE BRUSHED BRASS 1      LG CASE  
3      green antique    MFGR#3 MFGR#32 MFGR#3210      dark      STANDARD POLISHED BRASS 21      WRAP CASE  
4      metallic smoke   MFGR#1 MFGR#14 MFGR#1426      chocolate    SMALL PLATED BRASS 14      MED DRUM  
5      blush chiffon    MFGR#4 MFGR#45 MFGR#4510      forest      STANDARD POLISHED TIN 15      SM PKG  
6      ivory azure      MFGR#2 MFGR#23 MFGR#2325      white      PROMO PLATED STEEL 4      MED BAG  
7      blanched tan     MFGR#5 MFGR#51 MFGR#513      blue      SMALL PLATED COPPER 45      SM BAG  
8      khaki cream      MFGR#1 MFGR#13 MFGR#1328      ivory      PROMO BURNISHED TIN 41      LG DRUM  
9      rose moccasin     MFGR#4 MFGR#41 MFGR#4117      thistle     SMALL BURNISHED STEEL 12      WRAP CASE  
10     moccasin royal    MFGR#2 MFGR#21 MFGR#2128      floral      LARGE BURNISHED STEEL 44      LG CAN
```

```
Downloads — ec2-user@ip-172-31-28-103:~/apache-hive-2.0.1-bin —
GNU nano 2.9.8

#!/usr/bin/python
import sys

for line in sys.stdin:
    line = line.strip()
    vals = line.split('|')
    name = vals[1]
    namesplit = name.split(' ')
    newname= namesplit[1]+'_'+namesplit[0]
    vals[1] = newname
    print '\t'.join(vals)
```

Use Hive user defined function (i.e., SELECT TRANSFORM from our example, weekday mapper is available here: [http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/weekday\\_mapper.py](http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/weekday_mapper.py)) to perform the following transformation on Part table (creating a new PartSwap table): in the 2<sup>nd</sup> column/p\_name swap the two words in the column and replace the space by \_. For example, rose moccasin would become moccasin\_rose or honeydew dim would be dim\_honeydew.

```
create table partswap (
  p_partkey  int,
  p_name     varchar(22),
  p_mfgr     varchar(6),
  p_category varchar(7),
  p_brand1   varchar(9),
  p_color    varchar(11),
  p_type     varchar(25),
  p_size     int,
  p_container varchar(10)
)
```

**ROW FORMAT DELIMITED FIELDS**

**TERMINATED BY '\t' STORED AS TEXTFILE:**

**ADD FILE /home/ec2-user/apache-hive-2.0.1-bin/namemapperr.py:**

```
SELECT TRANSFORM (p_partkey, p_name, p_mfgr, p_category, p_brand1, p_color, p_type, p_size,
p_container) USING 'python namemapperr.py' AS (p_partkey, p_name, p_mfgr, p_category, p_brand1,
p_color, p_type, p_size, p_container) FROM part;
```

```
select * from partswap;
```

```
1  spring_lace      MFGR#1  MFGR#11 MFGR#1121  goldenrod  PROMO BURNISHED COPPER 7      JUMBO PKG
2  metallic_rosy    MFGR#4  MFGR#43 MFGR#4318  blush     LARGE BRUSHED BRASS   1      LG CASE
3  antique_green    MFGR#3  MFGR#32 MFGR#3210  dark      STANDARD POLISHED BRASS 21     WRAP CASE
4  smoke_metallic    MFGR#1  MFGR#14 MFGR#1426  chocolate  SMALL PLATED BRASS    14     MED DRUM
5  chiffon_blush    MFGR#4  MFGR#45 MFGR#4510  forest     STANDARD POLISHED TIN   15     SM PKG
6  azure_ivory       MFGR#2  MFGR#23 MFGR#2325  white      PROMO PLATED STEEL     4      MED BAG
7  tan_blanched     MFGR#5  MFGR#51 MFGR#513   blue       SMALL PLATED COPPER    45     SM BAG
8  cream_khaki       MFGR#1  MFGR#13 MFGR#1328  ivory      PROMO BURNISHED TIN    41     LG DRUM
9  moccasin_rose     MFGR#4  MFGR#41 MFGR#4117  thistle    SMALL BURNISHED STEEL  12     WRAP CASE
10 royal_moccasin    MFGR#2  MFGR#21 MFGR#2128  floral     LARGE BURNISHED STEEL  44     LG CAN
11 sandy_turquoise   MFGR#3  MFGR#34 MFGR#3438  chocolate  STANDARD BURNISHED NICKEL 43     WRAP BOX
12 olive_ivory       MFGR#3  MFGR#35 MFGR#3524  peru       MEDIUM ANODIZED STEEL  25     JUMBO CASE
13 olive_blue        MFGR#5  MFGR#53 MFGR#5333  ghost      MEDIUM BURNISHED NICKEL 1      JUMBO PACK
```

Keep in mind that your transform python code (split/join) should always use tab ('\t') between fields even if the source data is |-separated. You can also take a look at the transform example included with this assignment for your reference (Examples\_Assignment3.doc)

5) Download and install Pig:

```
cd
wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/pig-0.15.0.tar.gz
gunzip pig-0.15.0.tar.gz
tar xvf pig-0.15.0.tar
```

set the environment variables (this can also be placed in ~/.bashrc to make it permanent)

```
export PIG_HOME=/home/ec2-user/pig-0.15.0
export PATH=$PATH:$PIG_HOME/bin
```

Use the same vehicles file. Copy the vehicles.csv file to the HDFS if it is not already there.

Now run pig (and use the pig home variable we set earlier):

```
cd $PIG_HOME
bin/pig
```

Create the same table as what we used in Hive, assuming that vehicles.csv is in the home directory on HDFS:

```
hadoop dfs -copyFromLocal /home/ec2-user/vehicles.csv
/user/ec2-user
```

```
VehicleData = LOAD '/user/ec2-user/vehicles.csv' USING PigStorage(',')
AS (barrels08:FLOAT, barrelsA08:FLOAT, charge120:FLOAT, charge240:FLOAT, city08:FLOAT);
```

You can see the table description by

**DESCRIBE VehicleData;**

Verify that your data has loaded by running:

**VehicleG = GROUP VehicleData ALL;**

**Count = FOREACH VehicleG GENERATE COUNT(VehicleData);**

**DUMP Count;**

How many rows did you get? (if you get an error here, it is likely because vehicles.csv is not in HDFS)

```
Success!
Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReducetime  Alias  Featu
re Outputs
job_1581455066244_0024  1  1  5  5  5  5  4  4  4  4  Count,VehicleData,VehicleG  GROUP_BY,COMBINER  hdfs://localh
ost/tmp/temp140918457/tmp1223200921.

Input(s):
Successfully read 34175 records (11766951 bytes) from: "/user/ec2-user/vehicles.csv"

Output(s):
Successfully stored 1 records (9 bytes) in: "hdfs://localhost/tmp/temp140918457/tmp1223200921"

Counters:
Total records written : 1
Total bytes written : 9
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0
```

Create the same ThreeColExtract file that you have in the previous assignment, by placing barrels08, city08 and charge120 into a new file using PigStorage .You want the STORE command to record output in HDFS. (discussed in p457, Pig Chapter, “Data Processing Operator section)

NOTE: You can use this to get one column:

**OneCol = FOREACH VehicleData GENERATE barrels08, city08, charge120;**

**Store OneCol INTO ThreeColExtract1 USING PigStorage(',');**

Verify that the new file has been created and report the size of the newly created file. (you can use **quit** to exit the grunt shell)

```
Success!
Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReducetime  Alias  Fea
ture  Outputs
job_1581455066244_0036  1  0  5  5  5  5  0  0  0  0  VehicleData,a  MAP_ONLY  hdfs://localhost/user/ec2-user/Thre
eColExtract1.

Input(s):
Successfully read 34175 records (11766951 bytes) from: "/user/ec2-user/vehicles.csv"

Output(s):
Successfully stored 34175 records (627867 bytes) in: "hdfs://localhost/user/ec2-user/ThreeColExtract1"

Counters:
Total records written : 34175
Total bytes written : 627867
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1581455066244_0036

[ec2-user@ip-172-31-28-103 apache-hive-2.0.1-bin]$ hadoop fs -ls /user/ec2-user
Found 5 items
drwxr-xr-x - ec2-user supergroup 0 2020-02-12 01:46 /user/ec2-user/ThreeColExtract
drwxr-xr-x - ec2-user supergroup 0 2020-02-12 02:32 /user/ec2-user/ThreeColExtract1
drwxr-xr-x - ec2-user supergroup 0 2020-02-12 02:20 /user/ec2-user/ThreeColExtract_new
drwxr-xr-x - ec2-user supergroup 0 2020-02-12 01:36 /user/ec2-user/data
-rw-r--r-- 1 ec2-user supergroup 11766581 2020-02-12 01:39 /user/ec2-user/vehicles.csv
[ec2-user@ip-172-31-28-103 apache-hive-2.0.1-bin]$ hadoop fs -ls /user/ec2-user/ThreeColExtract1
Found 2 items
-rw-r--r-- 1 ec2-user supergroup 0 2020-02-12 02:32 /user/ec2-user/ThreeColExtract1/_SUCCESS
-rw-r--r-- 1 ec2-user supergroup 627867 2020-02-12 02:32 /user/ec2-user/ThreeColExtract1/part-m-000000
```

Submit a single document containing your written answers. Be sure that this document contains your name and “CSC 555 Assignment 3” at the top.